

final_proj

May 26, 2020

1 The Covid-19 Pandemic: How Healthcare and Economic Factors Influence Death Rate

Eshani Mehta, Emily Wang, Shauna Hannani

1.0.1 Importing Libraries

We start by importing the libraries we will use later in the project.

```
[1]: import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
%matplotlib inline

import json
from urllib.request import urlopen

import seaborn as sns
sns.set(style = "whitegrid",
        color_codes = True,
        font_scale = 1.5)

import plotly.express as px
import plotly.graph_objects as go

from sklearn import linear_model as lm
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
from sklearn import preprocessing
```

1.0.2 Reading in Data

We then read the datasets that we will use throughout this project. * The counties dataset contains county-based information such as health factors, population, and some economic factors. * The

confirmed dataset contains the number of cases in each county through 5/12/20. * The deaths dataset contains the number of deaths in each county through 5/12/20.

We also brought in 2 extra datasets from the USDA's county level data sets (link: <https://www.ers.usda.gov/data-products/county-level-data-sets/download-data/>). * The poverty dataset contains county-based information such as the percent of people under poverty level. * The unemployment dataset contains county-based information such as the unemployment rate in 2018 and the median household income in 2018.

```
[2]: #states = pd.read_csv('finalproj/data/5.12states.csv')
      counties = pd.read_csv('abridged_couties.csv')
      confirmed = pd.read_csv('new_time_series_covid19_confirmed_US.csv')
      deaths = pd.read_csv('new_time_series_covid19_deaths_US.csv')

      #Following data sets taken from the USDA's county level data sets
      poverty = pd.read_csv('PovertyEstimates.csv')
      unemployment = pd.read_csv('Unemployment.csv', thousands=',')
```

```
[3]: counties.head()
```

```
[3]:  countyFIPS  STATEFP  COUNTYFP  CountyName  StateName  State      lat  \
0      01001      1.0      1.0    Autauga        AL    Alabama  32.540091
1      01003      1.0      3.0    Baldwin        AL    Alabama  30.738314
2      01005      1.0      5.0    Barbour        AL    Alabama  31.874030
3      01007      1.0      7.0      Bibb        AL    Alabama  32.999024
4      01009      1.0      9.0    Blount        AL    Alabama  33.990440

      lon  POP_LATITUDE  POP_LONGITUDE  ... >500 gatherings public schools  \
0 -86.645649      32.500389      -86.494165  ...      737497.0      737500.0
1 -87.726272      30.548923      -87.762381  ...      737497.0      737500.0
2 -85.397327      31.844036      -85.310038  ...      737497.0      737500.0
3 -87.125260      33.030921      -87.127659  ...      737497.0      737500.0
4 -86.562711      33.955243      -86.591491  ...      737497.0      737500.0

      restaurant dine-in  entertainment/gym  federal guidelines  \
0      737503.0      737512.0      737500.0
1      737503.0      737512.0      737500.0
2      737503.0      737512.0      737500.0
3      737503.0      737512.0      737500.0
4      737503.0      737512.0      737500.0

      foreign travel ban  SVIPercentile  HPSAShortage  HPSAServedPop  \
0      737495.0      0.4354      NaN      NaN
1      737495.0      0.2162      NaN      NaN
2      737495.0      0.9959      6.08      5400.0
3      737495.0      0.6003      2.75      14980.0
4      737495.0      0.4242      7.21      31850.0
```

```

      HPSAUnderservedPop
0          NaN
1          NaN
2       18241.0
3        6120.0
4       25233.0

```

[5 rows x 87 columns]

```
[4]: confirmed.head()
```

```

[4]:   UID iso2 iso3  code3  FIPS Admin2      Province_State Country_Region \
0    16  AS  ASM    16  60.0    NaN      American Samoa          US
1   316  GU  GUM   316  66.0    NaN              Guam          US
2   580  MP  MNP   580  69.0    NaN  Northern Mariana Islands      US
3   630  PR  PRI   630  72.0    NaN      Puerto Rico          US
4   850  VI  VIR   850  78.0    NaN      Virgin Islands          US

```

```

      Lat      Long_ ... 5/3/20 5/4/20 5/5/20 5/6/20 5/7/20 5/8/20 \
0 -14.2710 -170.1320 ...    0      0      0      0      0      0
1  13.4443  144.7937 ...   145    145    145    149    149    151
2  15.0979  145.6739 ...    14     14     14     15     15     15
3  18.2208  -66.5901 ...  1808   1843   1924   1968   2031   2156
4  18.3358  -64.8963 ...    66     66     66     66     66     68

```

```

      5/9/20 5/10/20 5/11/20 5/12/20
0         0         0         0         0
1        151        151        151        152
2         16         16         19         19
3        2173        2198        2256        2299
4         68         69         69         69

```

[5 rows x 123 columns]

```
[5]: deaths.head()
```

```

[5]:   UID iso2 iso3  code3  FIPS Admin2      Province_State Country_Region \
0    16  AS  ASM    16  60.0    NaN      American Samoa          US
1   316  GU  GUM   316  66.0    NaN              Guam          US
2   580  MP  MNP   580  69.0    NaN  Northern Mariana Islands      US
3   630  PR  PRI   630  72.0    NaN      Puerto Rico          US
4   850  VI  VIR   850  78.0    NaN      Virgin Islands          US

```

```

      Lat      Long_ ... 5/3/20 5/4/20 5/5/20 5/6/20 5/7/20 5/8/20 \
0 -14.2710 -170.1320 ...    0      0      0      0      0      0
1  13.4443  144.7937 ...     5      5      5      5      5      5
2  15.0979  145.6739 ...     2      2      2      2      2      2

```

| | | | | | | | | | |
|---|---------|----------|-----|----|----|----|----|-----|-----|
| 3 | 18.2208 | -66.5901 | ... | 97 | 97 | 99 | 99 | 102 | 107 |
| 4 | 18.3358 | -64.8963 | ... | 4 | 4 | 4 | 4 | 4 | 4 |

| | 5/9/20 | 5/10/20 | 5/11/20 | 5/12/20 |
|---|--------|---------|---------|---------|
| 0 | 0 | 0 | 0 | 0 |
| 1 | 5 | 5 | 5 | 5 |
| 2 | 2 | 2 | 2 | 2 |
| 3 | 108 | 111 | 113 | 114 |
| 4 | 4 | 4 | 5 | 6 |

[5 rows x 124 columns]

```
[6]: poverty.head()
```

```
[6]: FIPStxt Stabr Area_name Rural-urban_Continuum_Code_2003 \
0      0      US United States NaN
1    1000      AL Alabama NaN
2    1001      AL Autauga County 2.0
3    1003      AL Baldwin County 4.0
4    1005      AL Barbour County 6.0
```

| | Urban_Influence_Code_2003 | Rural-urban_Continuum_Code_2013 | \ |
|---|---------------------------|---------------------------------|---|
| 0 | NaN | NaN | |
| 1 | NaN | NaN | |
| 2 | 2.0 | 2.0 | |
| 3 | 5.0 | 3.0 | |
| 4 | 6.0 | 6.0 | |

| | Urban_Influence_Code_2013 | POVALL_2018 | CI90LBAll_2018 | CI90UBAll_2018 | ... | \ |
|---|---------------------------|-------------|----------------|----------------|-----|---|
| 0 | NaN | 41,852,315 | 41,619,366 | 42,085,264 | ... | |
| 1 | NaN | 801,758 | 785,668 | 817,848 | ... | |
| 2 | 2.0 | 7,587 | 6,334 | 8,840 | ... | |
| 3 | 2.0 | 21,069 | 17,390 | 24,748 | ... | |
| 4 | 6.0 | 6,788 | 5,662 | 7,914 | ... | |

| | CI90UB517P_2018 | MEDHHINC_2018 | CI90LBINC_2018 | CI90UBINC_2018 | POV04_2018 | \ |
|---|-----------------|---------------|----------------|----------------|------------|---|
| 0 | 17.2 | 61,937 | 61,843 | 62,031 | 3,758,704 | |
| 1 | 23.7 | 49,881 | 49,123 | 50,639 | 73,915 | |
| 2 | 23.9 | 59,338 | 53,628 | 65,048 | NaN | |
| 3 | 16.9 | 57,588 | 54,437 | 60,739 | NaN | |
| 4 | 45.9 | 34,382 | 31,157 | 37,607 | NaN | |

| | CI90LB04_2018 | CI90UB04_2018 | PCTPOV04_2018 | CI90LB04P_2018 | CI90UB04P_2018 |
|---|---------------|---------------|---------------|----------------|----------------|
| 0 | 3,714,862 | 3,802,546 | 19.5 | 19.3 | 19.7 |
| 1 | 69,990 | 77,840 | 26.0 | 24.6 | 27.4 |
| 2 | NaN | NaN | NaN | NaN | NaN |
| 3 | NaN | NaN | NaN | NaN | NaN |

| | | | | | |
|---|-----|-----|-----|-----|-----|
| 4 | NaN | NaN | NaN | NaN | NaN |
|---|-----|-----|-----|-----|-----|

[5 rows x 34 columns]

```
[7]: unemployment.head()
```

```
[7]:
```

| | FIPStxt | Stabr | area_name | Rural_urban_continuum_code_2013 | \ |
|---|---------|-------|--------------------|---------------------------------|---|
| 0 | 0 | US | United States | NaN | |
| 1 | 1000 | AL | Alabama | NaN | |
| 2 | 1001 | AL | Autauga County, AL | 2.0 | |
| 3 | 1003 | AL | Baldwin County, AL | 3.0 | |
| 4 | 1005 | AL | Barbour County, AL | 6.0 | |

| | Urban_influence_code_2013 | Metro_2013 | Civilian_labor_force_2000 | \ |
|---|---------------------------|------------|---------------------------|---|
| 0 | NaN | NaN | 142601667.0 | |
| 1 | NaN | NaN | 2133223.0 | |
| 2 | 2.0 | 1.0 | 21720.0 | |
| 3 | 2.0 | 1.0 | 69533.0 | |
| 4 | 6.0 | 0.0 | 11373.0 | |

| | Employed_2000 | Unemployed_2000 | Unemployment_rate_2000 | ... | \ |
|---|---------------|-----------------|------------------------|-----|---|
| 0 | 136904680.0 | 5696987.0 | 4.0 | ... | |
| 1 | 2035594.0 | 97629.0 | 4.6 | ... | |
| 2 | 20846.0 | 874.0 | 4.0 | ... | |
| 3 | 66971.0 | 2562.0 | 3.7 | ... | |
| 4 | 10748.0 | 625.0 | 5.5 | ... | |

| | Civilian_labor_force_2018 | Employed_2018 | Unemployed_2018 | \ |
|---|---------------------------|---------------|-----------------|---|
| 0 | 161389026.0 | 155102319.0 | 6286707.0 | |
| 1 | 2216627.0 | 2130845.0 | 85782.0 | |
| 2 | 26196.0 | 25261.0 | 935.0 | |
| 3 | 95233.0 | 91809.0 | 3424.0 | |
| 4 | 8414.0 | 7987.0 | 427.0 | |

| | Unemployment_rate_2018 | Civilian_labor_force_2019 | Employed_2019 | \ |
|---|------------------------|---------------------------|---------------|---|
| 0 | 3.9 | 163100055.0 | 157115247.0 | |
| 1 | 3.9 | 2241747.0 | 2174483.0 | |
| 2 | 3.6 | 26172.0 | 25458.0 | |
| 3 | 3.6 | 97328.0 | 94675.0 | |
| 4 | 5.1 | 8537.0 | 8213.0 | |

| | Unemployed_2019 | Unemployment_rate_2019 | Median_Household_Income_2018 | \ |
|---|-----------------|------------------------|------------------------------|---|
| 0 | 5984808.0 | 3.7 | 61937.0 | |
| 1 | 67264.0 | 3.0 | 49881.0 | |
| 2 | 714.0 | 2.7 | 59338.0 | |
| 3 | 2653.0 | 2.7 | 57588.0 | |
| 4 | 324.0 | 3.8 | 34382.0 | |

| | Med_HH_Income_Percent_of_State_Total_2018 |
|---|---|
| 0 | NaN |
| 1 | 100.0 |
| 2 | 119.0 |
| 3 | 115.5 |
| 4 | 68.9 |

[5 rows x 88 columns]

1.1 Data Cleaning

1.1.1 Merging Relevant Datasets

The first step is creating a merged dataset with each county’s information, the number of confirmed cases, the number of deaths, and other relevant information from other datasets.

We create a copy to ensure that the original `counties` dataframe is intact.

Since the “countyFIPS” column currently is type str, we convert it to numbers so we can match it with the “FIPS” columns in `confirmed` and `deaths`. We also drop any empty values for FIPS since all valid counties have a FIPS value.

The FIPS columns serve as a primary key since it’s a standardized id to identify counties that is present in most county datasets even online.

We first inner merge `merged` with `confirmed` on the FIPS columns, specifically only looking at the “5/12/20” column since that contains the most updated number of cases per county. We then rename this to a more usable name so it’ll be easier to reference later.

We then similarly inner merge the new `merged` with `deaths` on the FIPS columns, specifically only looking at the “5/12/20” column since that contains the most updated number of deaths per county. We then rename this to a more usable name so it’ll be easier to reference later.

```
[8]: merged = counties.copy()

merged["countyFIPS"] = pd.to_numeric(merged["countyFIPS"], errors='coerce')
merged = merged.dropna(subset=["countyFIPS"])

#merge w confirmed
merged = merged.merge(right=confirmed[["FIPS", "5/12/20"]], how='inner',
    ↳left_on='countyFIPS', right_on='FIPS')
merged = merged.rename(columns={'5/12/20' : 'confirmed'})

#merged w deaths
merged = merged.merge(right=deaths[["FIPS", "5/12/20"]], how='inner',
    ↳left_on='countyFIPS', right_on='FIPS')
merged = merged.rename(columns={'5/12/20' : 'deaths'})
```

```
#drop irrelevant columns
merged = merged.drop(["FIPS_x", "FIPS_y"], axis = 1)
merged
```

```
[8]:
```

| | countyFIPS | STATEFP | COUNTYFP | CountyName \ |
|------|------------|---------|----------|-----------------------------------|
| 0 | 1001.0 | 1.0 | 1.0 | Autauga |
| 1 | 1003.0 | 1.0 | 3.0 | Baldwin |
| 2 | 1005.0 | 1.0 | 5.0 | Barbour |
| 3 | 1007.0 | 1.0 | 7.0 | Bibb |
| 4 | 1009.0 | 1.0 | 9.0 | Blount |
| ... | ... | ... | ... | ... |
| 3135 | 2195.0 | 2.0 | 195.0 | Petersburg Borough |
| 3136 | 2198.0 | 2.0 | 198.0 | Prince of Wales-Hyder Census Area |
| 3137 | 2230.0 | 2.0 | 230.0 | Skagway Municipality |
| 3138 | 2275.0 | 2.0 | 275.0 | Wrangell City and Borough |
| 3139 | 15005.0 | 15.0 | 5.0 | Kalawao |

| | StateName | State | lat | lon | POP_LATITUDE | POP_LONGITUDE \ |
|------|-----------|---------|-----------|------------|--------------|-----------------|
| 0 | AL | Alabama | 32.540091 | -86.645649 | 32.500389 | -86.494165 |
| 1 | AL | Alabama | 30.738314 | -87.726272 | 30.548923 | -87.762381 |
| 2 | AL | Alabama | 31.874030 | -85.397327 | 31.844036 | -85.310038 |
| 3 | AL | Alabama | 32.999024 | -87.125260 | 33.030921 | -87.127659 |
| 4 | AL | Alabama | 33.990440 | -86.562711 | 33.955243 | -86.591491 |
| ... | ... | ... | ... | ... | ... | ... |
| 3135 | AK | NaN | NaN | NaN | 56.812712 | -133.115025 |
| 3136 | AK | NaN | NaN | NaN | 55.448164 | -132.560842 |
| 3137 | AK | NaN | NaN | NaN | 59.464536 | -135.311501 |
| 3138 | AK | NaN | NaN | NaN | 56.385821 | -132.310837 |
| 3139 | HI | NaN | NaN | NaN | 21.188495 | -156.979972 |

| | ... | restaurant dine-in | entertainment/gym | federal guidelines \ |
|------|-----|--------------------|-------------------|----------------------|
| 0 | ... | 737503.0 | 737512.0 | 737500.0 |
| 1 | ... | 737503.0 | 737512.0 | 737500.0 |
| 2 | ... | 737503.0 | 737512.0 | 737500.0 |
| 3 | ... | 737503.0 | 737512.0 | 737500.0 |
| 4 | ... | 737503.0 | 737512.0 | 737500.0 |
| ... | ... | ... | ... | ... |
| 3135 | ... | 737501.0 | 737501.0 | 737500.0 |
| 3136 | ... | 737501.0 | 737501.0 | 737500.0 |
| 3137 | ... | 737501.0 | 737501.0 | 737500.0 |
| 3138 | ... | 737501.0 | 737501.0 | 737500.0 |
| 3139 | ... | 737504.0 | 737509.0 | 737500.0 |

| | foreign travel ban | SVIPercentile | HPSAShortage | HPSAServedPop \ |
|---|--------------------|---------------|--------------|-----------------|
| 0 | 737495.0 | 0.4354 | NaN | NaN |
| 1 | 737495.0 | 0.2162 | NaN | NaN |

| | | | | |
|------|----------|--------|------|---------|
| 2 | 737495.0 | 0.9959 | 6.08 | 5400.0 |
| 3 | 737495.0 | 0.6003 | 2.75 | 14980.0 |
| 4 | 737495.0 | 0.4242 | 7.21 | 31850.0 |
| ... | ... | ... | ... | ... |
| 3135 | 737495.0 | 0.6650 | NaN | NaN |
| 3136 | 737495.0 | 0.7662 | 1.50 | 1050.0 |
| 3137 | 737495.0 | 0.1685 | 0.69 | 175.0 |
| 3138 | 737495.0 | 0.5618 | NaN | NaN |
| 3139 | 737495.0 | 0.3162 | NaN | NaN |

| | HPSAUnderservedPop | confirmed | deaths |
|------|--------------------|-----------|--------|
| 0 | NaN | 91 | 4 |
| 1 | NaN | 227 | 7 |
| 2 | 18241.0 | 67 | 1 |
| 3 | 6120.0 | 46 | 1 |
| 4 | 25233.0 | 45 | 0 |
| ... | ... | ... | ... |
| 3135 | NaN | 4 | 1 |
| 3136 | 5260.0 | 2 | 0 |
| 3137 | 2412.0 | 0 | 0 |
| 3138 | NaN | 0 | 0 |
| 3139 | NaN | 0 | 0 |

[3140 rows x 89 columns]

1.1.2 Merging External Datasets

Here, we merge the two USDA datasets in a similar fashion to what we did above. We convert the FIPS columns in `poverty` and `unemployment` to numeric data again and select the columns from each that we want to use in our analysis before performing an inner merge.

```
[9]: merged2 = merged.copy()
poverty["FIPStxt"] = pd.to_numeric(poverty["FIPStxt"], errors='coerce')
unemployment["FIPStxt"] = pd.to_numeric(unemployment["FIPStxt"],
    ↪errors='coerce')

#merge w poverty to add and rename pov_pct
merged2 = merged2.merge(right=poverty[["FIPStxt", "PCTPOVALL_2018"]],
    ↪how='inner', left_on='countyFIPS', right_on='FIPStxt')
merged2 = merged2.rename(columns={"PCTPOVALL_2018" : "pov_pct"})

#merge w unemployment to add and rename unemploy_rate and
    ↪median_household_income
merged2 = merged2.merge(right=unemployment[["FIPStxt",
    ↪"Unemployment_rate_2018", "Median_Household_Income_2018"]], how='inner',
    ↪left_on='countyFIPS', right_on='FIPStxt')
```



```
merged2 = merged2.rename(columns={"Unemployment_rate_2018" : "unemploy_rate",
↳ "Median_Household_Income_2018": "med_income"})
```

1.1.3 Adding Relevant Columns and Some Data Cleaning

We create a new copy of all the merged datasets. We want to first add a `case_rate` column, which is just the # of cases divided by the total population of a county. We also multiply this by 100 to make it easier to read.

We want to then add a `death_rate` column, which is the # of deaths divided by the # of confirmed cases. Before we this, we remove the 272 counties that have 0 confirmed cases so we don't have a division by 0 error. Since this project is also aiming to look at effects of certain factors on the `case_rate` and `death_rate`, counties with no cases don't add anything to our data. We also multiply this by 100 to make it easier to read.

Other columns we added: * `old` is the percent of the population that is age 65+ * `inMedicare` is the percent of Medicare eligible people who are actually enrolled * `medicare_rate` is the percent of the population that is eligible for Medicare

We also noticed in the dataset that counties with latitude and longitude equal to 0 or missing were other territories or not valid counties or lacking important information, so we remove all instances where that is true.

We decided we should analyze which columns were missing more data than others, which ended up including include ["3-Yr Diabetes", all the "3-YrMortality", "stay at home" (396), "HPSA" (1013)].

We also checked on what states were included in the data after cleaning. The data now covered 48 states (excluding Alaska and Hawaii since they're not continental and including District of Columbia).

```
[10]: #add case rate
data = merged2.copy()
data["case_rate"] = data["confirmed"] / data["PopulationEstimate2018"]
data["case_rate"] = data["case_rate"] * 100

#add death rate
data = data.loc[data["confirmed"] != 0]
data["death_rate"] = data["deaths"] / data["confirmed"]
data["death_rate"] = data["death_rate"] * 100

#add age column
data["old"] = (data["PopulationEstimate65+2017"] /
↳ data["PopulationEstimate2018"]) * 100

#add medicare cols
data = data.dropna(subset=["MedicareEnrollment,AgedTot2017"])
data["inMedicare"] = (data["MedicareEnrollment,AgedTot2017"] /
↳ data["#EligibleforMedicare2018"]) * 100
```

```

data["medicare_rate"] = (data["#EligibleforMedicare2018"] /
↳data["PopulationEstimate2018"]) * 100

#add hospitals and icu_beds per 1000 people
data["hospitals/1000ppl"] = (data["#Hospitals"] /
↳data["PopulationEstimate2018"]) * 1000
data["icu_beds/1000ppl"] = (data["#ICU_beds"] / data["PopulationEstimate2018"])*
↳* 1000

#drop null or 0 latitudes
data = data.loc[data["lat"] != 0]
data = data.dropna(subset=["lat"])

data.isna().sum()
#note: columns with many na values include ["3-Yr Diabetes", all the
↳"3-YrMortality", "stay at home" (396), "HPSA" (1013)]

data.State.unique()
#48 states (excluding Alaska and Hawaii since they're not continental and
↳including District of Columbia)

data

```

```

[10]:
   countyFIPS  STATEFP  COUNTYFP  CountyName  StateName  State \
0      1001.0      1.0      1.0    Autauga        AL    Alabama
1      1003.0      1.0      3.0    Baldwin        AL    Alabama
2      1005.0      1.0      5.0    Barbour        AL    Alabama
3      1007.0      1.0      7.0      Bibb        AL    Alabama
4      1009.0      1.0      9.0    Blount        AL    Alabama
...      ...      ...      ...      ...      ...      ...
3128    56039.0     56.0     39.0      Teton        WY    Wyoming
3129    56041.0     56.0     41.0      Uinta        WY    Wyoming
3130    56043.0     56.0     43.0    Washakie        WY    Wyoming
3132     8014.0      8.0     14.0  Broomfield        CO    Colorado
3133    12086.0     12.0     86.0  Miami-Dade        FL    Florida

   lat      lon  POP_LATITUDE  POP_LONGITUDE  ...  FIPStxt_y \
0  32.540091 -86.645649    32.500389    -86.494165  ...    1001
1  30.738314 -87.726272    30.548923    -87.762381  ...    1003
2  31.874030 -85.397327    31.844036    -85.310038  ...    1005
3  32.999024 -87.125260    33.030921    -87.127659  ...    1007
4  33.990440 -86.562711    33.955243    -86.591491  ...    1009
...      ...      ...      ...      ...      ...      ...
3128  43.713556 -110.570974    43.494174    -110.784353  ...    56039
3129  41.289323 -110.553036    41.271860    -110.767519  ...    56041
3130  43.909060 -107.679282    44.012142    -107.911552  ...    56043
3132  39.963039 -105.058542    39.936888    -105.055491  ...     8014

```

| | | | | | | |
|------|-----------|------------|-----------|------------|-----|-------|
| 3133 | 25.607895 | -80.587502 | 25.774565 | -80.298888 | ... | 12086 |
|------|-----------|------------|-----------|------------|-----|-------|

| | unemploy_rate | med_income | case_rate | death_rate | old | inMedicare \ |
|------|---------------|------------|-----------|------------|-----------|--------------|
| 0 | 3.6 | 59338.0 | 0.163666 | 4.395604 | 15.093254 | 70.338316 |
| 1 | 3.6 | 57588.0 | 0.104118 | 3.083700 | 19.453541 | 76.812263 |
| 2 | 5.1 | 34382.0 | 0.269282 | 1.492537 | 19.119006 | 70.438557 |
| 3 | 3.9 | 46064.0 | 0.205357 | 2.173913 | 16.214286 | 66.088144 |
| 4 | 3.5 | 50412.0 | 0.077801 | 0.000000 | 17.895920 | 72.425047 |
| ... | ... | ... | ... | ... | ... | ... |
| 3128 | 2.9 | 99087.0 | 0.424592 | 0.000000 | 14.509770 | 83.505976 |
| 3129 | 4.2 | 63401.0 | 0.044337 | 0.000000 | 13.315927 | 72.092376 |
| 3130 | 4.1 | 55190.0 | 0.101458 | 0.000000 | 21.280913 | 82.449182 |
| 3132 | 2.8 | 96924.0 | 0.290181 | 9.950249 | 13.072603 | 73.026762 |
| 3133 | 3.5 | 52043.0 | 0.520897 | 3.510601 | 15.898321 | 84.000125 |

| | medicare_rate | hospitals/1000ppl | icu_beds/1000ppl |
|------|---------------|-------------------|------------------|
| 0 | 20.573371 | 0.017985 | 0.107912 |
| 1 | 24.834650 | 0.013760 | 0.233921 |
| 2 | 26.851815 | 0.040191 | 0.200957 |
| 3 | 22.892857 | 0.044643 | 0.000000 |
| 4 | 22.778354 | 0.017289 | 0.103734 |
| ... | ... | ... | ... |
| 3128 | 16.312118 | 0.043326 | 0.259954 |
| 3129 | 17.705306 | 0.049264 | 0.295581 |
| 3130 | 25.580216 | 0.126823 | 0.000000 |
| 3132 | 17.047079 | 0.014437 | 0.288738 |
| 3133 | 17.370847 | 0.006156 | 0.214732 |

[2812 rows x 101 columns]

1.2 Exploratory Data Analysis

1.2.1 Effects of the timing of Stay at Home orders on case_rate and death_rate

Since a big portion of the news today is focused on stay at home orders, we wanted to see if there was any obvious effect of when the stay at home orders were put in place and the rates of cases and deaths in a county.

After calculating a more readable number of days for the “stay at home” column, we made boxplots for each of case_rate and death_rate. We decided to hide outliers to make the plot more readable, and there were many outliers that fell way beyond the current graph axes.

```
[11]: dates = data.copy()

#737427 is the proleptic Gregorian ordinal of January 1, 2020
#we subtract that so we can see how many days into the new year the order was
↳put into place
```

```

dates["stayHome"] = dates["stay at home"] - 737427

#drop all counties where there is no stay at home order established since this_
↳is about timing
dates = dates.dropna(subset=["stayHome"])

print("Stay at home dates range (in days since Jan. 1, 2020): ", "March 17, 2020 (", min(dates["stayHome"]), " days) - April 5, 2020 (", max(dates["stayHome"]), " days)")

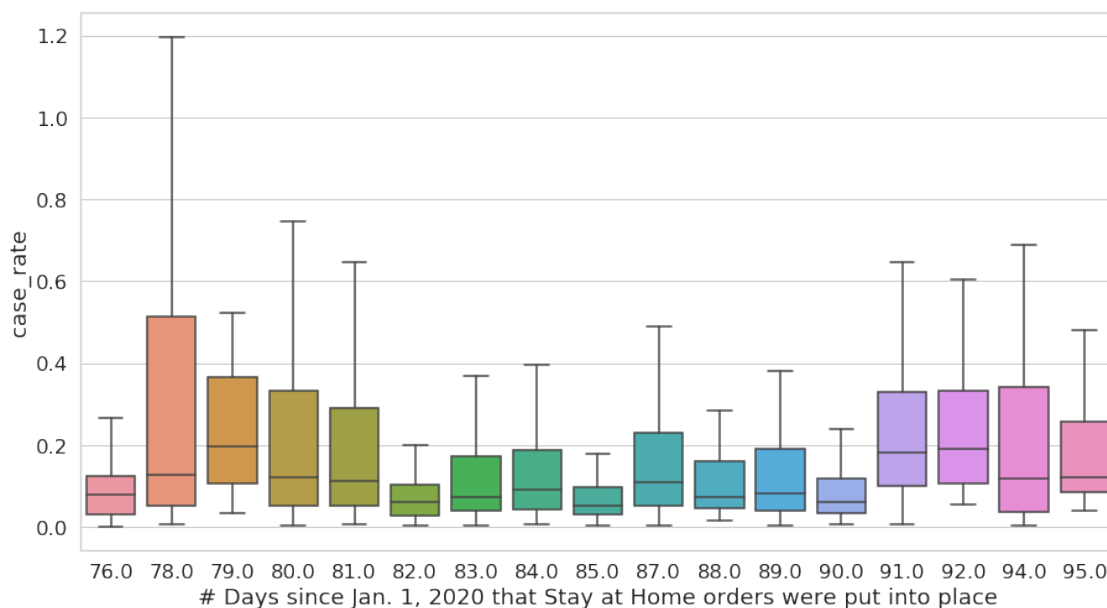
#boxplot of case_rate
plt.figure(figsize = (15, 8))
sns.boxplot(x="stayHome", y = "case_rate", data=dates, showfliers=False)
plt.xlabel("# Days since Jan. 1, 2020 that Stay at Home orders were put into place")

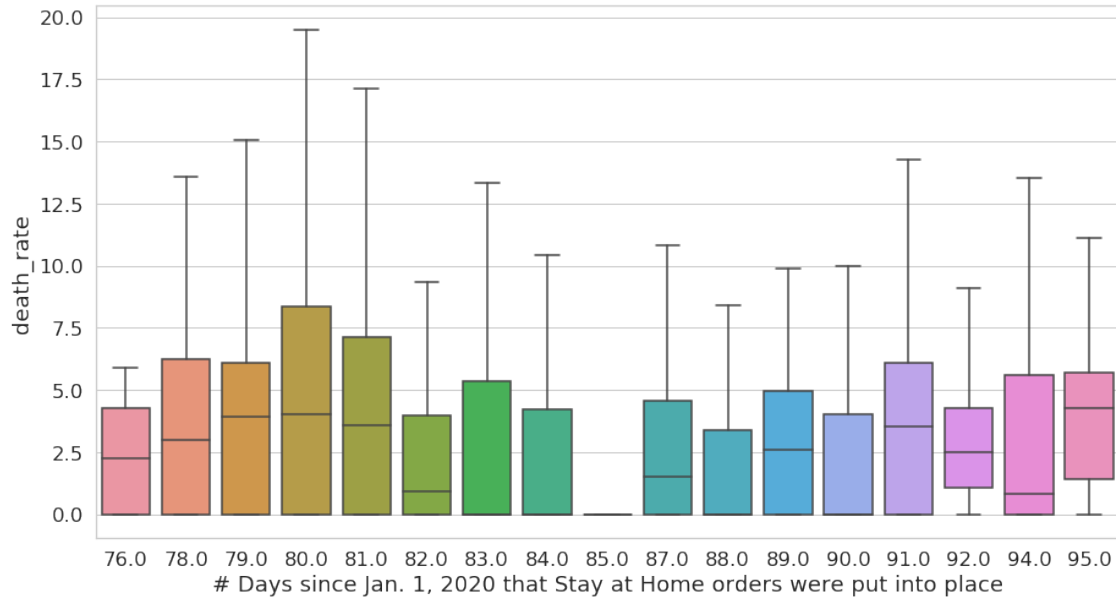
#boxplot of death_rate
plt.figure(figsize = (15, 8))
sns.boxplot(x="stayHome", y = "death_rate", data=dates, showfliers=False)
plt.xlabel("# Days since Jan. 1, 2020 that Stay at Home orders were put into place")

```

Stay at home dates range (in days since Jan. 1, 2020): March 17, 2020 (76.0 days) - April 5, 2020 (95.0 days)

[11]: Text(0.5, 0, '# Days since Jan. 1, 2020 that Stay at Home orders were put into place')





1.2.2 Political Affiliation of County and Correlation with case_rate and death_rate

We found the dem_to_rep_ratio column interesting, and since many people on either side have differing opinions about the pandemic and social distancing, we thought it would be interesting to see if there was any correlation to case_rate or death_rate.

We categorized the ratios into democratic and republican and then created boxplots for each factor.

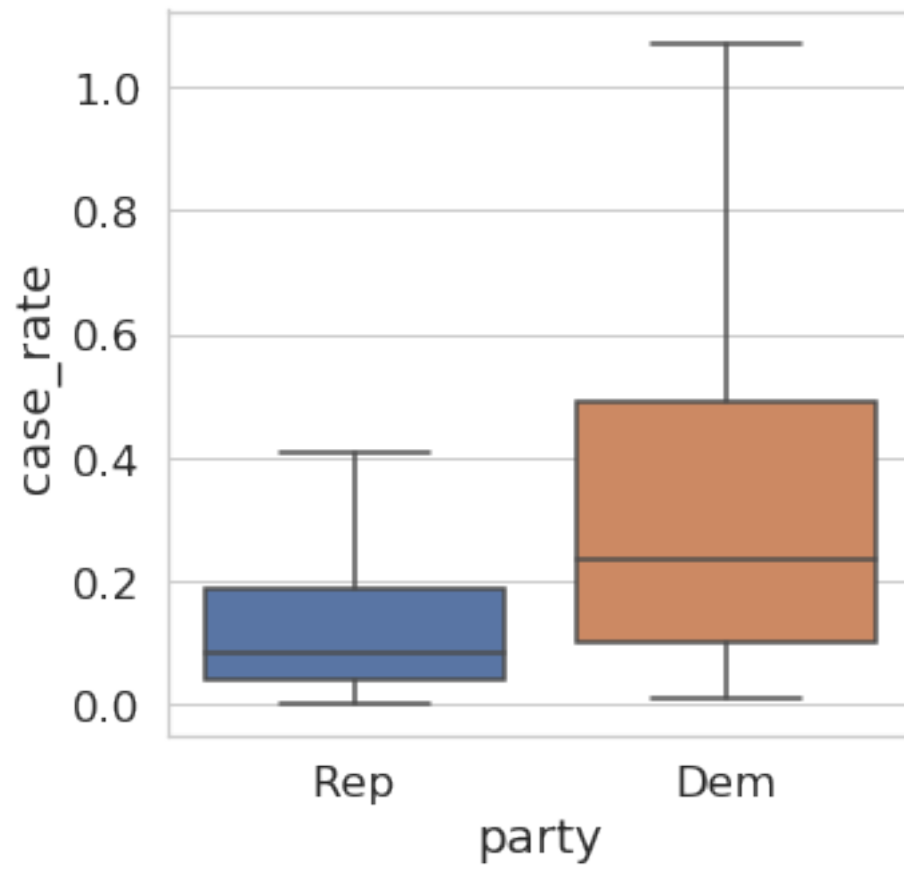
```
[12]: political = data.copy()
      political = political.dropna(subset=["dem_to_rep_ratio"])

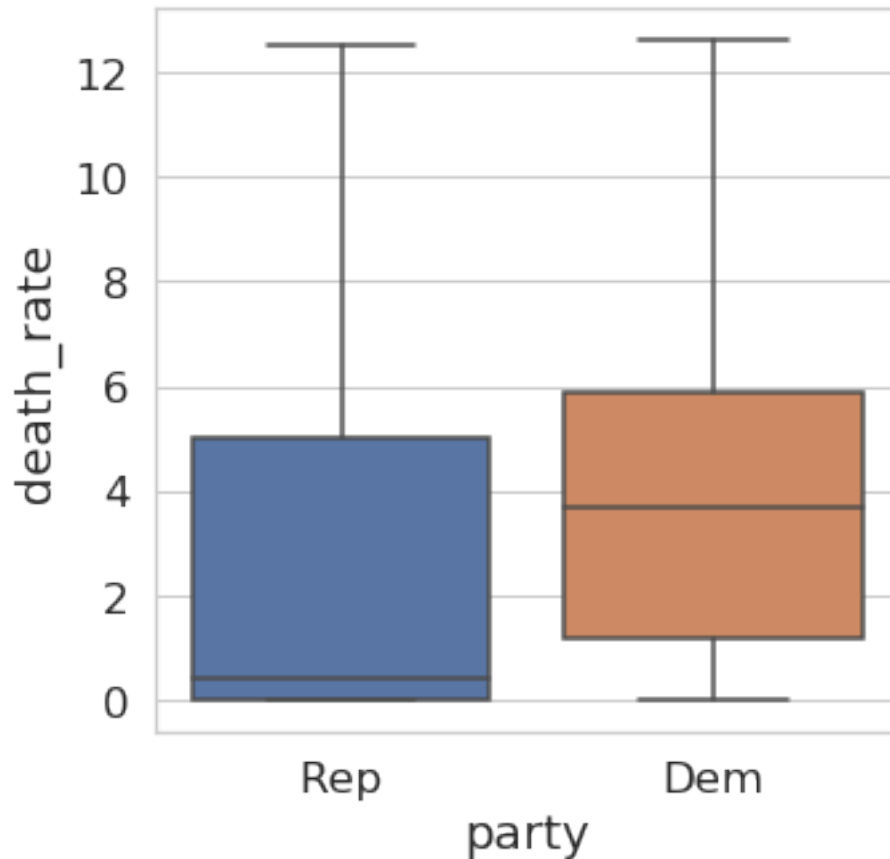
      #categorize
      political.loc[political["dem_to_rep_ratio"] > 1, "party"] = "Dem"
      political.loc[political["dem_to_rep_ratio"] < 1, "party"] = "Rep"
      political[["party"]]

      #make boxplot for case_rate
      plt.figure(figsize = (5, 5))
      sns.boxplot(x="party", y = "case_rate", data=political, showfliers=False)

      #make boxplot for death_rate
      plt.figure(figsize = (5, 5))
      sns.boxplot(x="party", y = "death_rate", data=political, showfliers=False)
```

```
[12]: <matplotlib.axes._subplots.AxesSubplot at 0x7f380c270710>
```





1.2.3 Exploring counties and the # confirmed cases

Here, we just wanted to know which counties had the most confirmed cases as of 5/12. We weren't surprised to see New York, of course, but many of the other names were unexpected.

We also calculated the mean and median number of confirmed cases per county. Since we already removed the 0 case counties, we were able to get a better value. The mean was obviously influenced by large outliers like the top 10 shown here. The median was interesting since we didn't expect it to be low because we usually only hear about the counties with large numbers of cases on the news. However, it makes sense that most counties don't have as many, especially ones in rural areas.

```
[13]: print("Mean number of confirmed cases per county: ", np.mean(data["confirmed"]))
      print("Median number of confirmed cases per county: ", np.
            ↪median(data["confirmed"]))
      print("Total number of counties: ", data.shape[0])

      #looking at the counties with the most cases and deaths
      most_cases = data.sort_values("confirmed", ascending = False)
      most_cases[["CountyName", "StateName", "confirmed"]].head(10)
```

Mean number of confirmed cases per county: 480.7528449502134
Median number of confirmed cases per county: 31.0
Total number of counties: 2812

```
[13]:
```

| | CountyName | StateName | confirmed |
|------|--------------|-----------|-----------|
| 1849 | New York | NY | 186123 |
| 601 | Cook | IL | 55470 |
| 1848 | Nassau | NY | 38434 |
| 1870 | Suffolk | NY | 37062 |
| 198 | Los Angeles | CA | 33211 |
| 1878 | Westchester | NY | 31472 |
| 2285 | Philadelphia | PA | 18537 |
| 1303 | Wayne | MI | 18274 |
| 1216 | Middlesex | MA | 17953 |
| 1773 | Hudson | NJ | 17677 |

1.2.4 Distribution of Case Rate

As we considered which factor to predict, we created some histograms to observe the distribution of `case_rate`.

```
[14]: hist_cases = data.copy()
hist_cases = hist_cases[["case_rate"]]

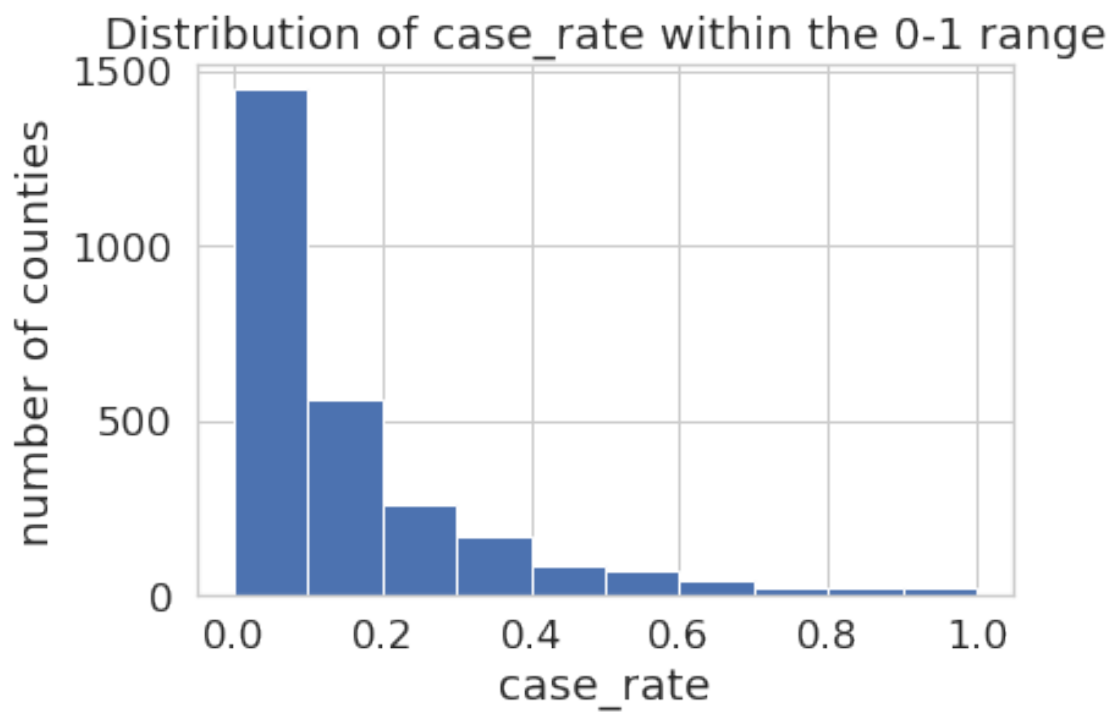
hist_cases.hist()
plt.title("Distribution of case_rate in all counties")
plt.xlabel("case_rate")
plt.ylabel("number of counties")

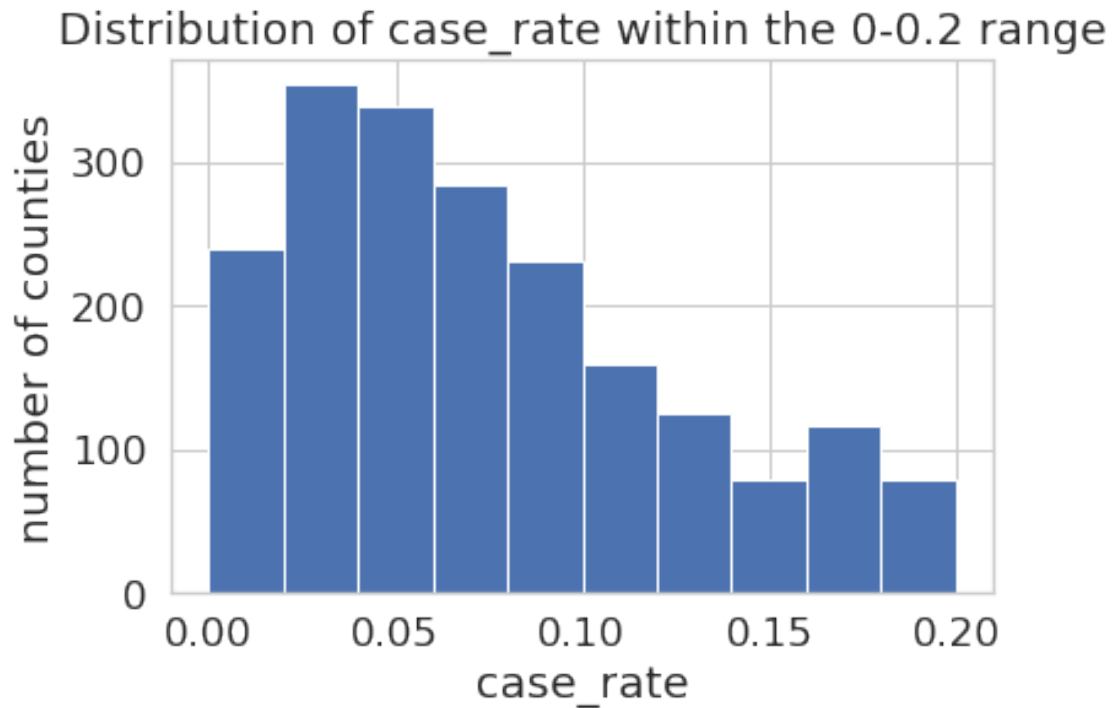
hist_cases.hist(range = (0,1))
plt.title("Distribution of case_rate within the 0-1 range")
plt.xlabel("case_rate")
plt.ylabel("number of counties")

hist_cases.hist(range = (0, 0.2))
plt.title("Distribution of case_rate within the 0-0.2 range")
plt.xlabel("case_rate")
plt.ylabel("number of counties")

hist_cases.loc[hist_cases["case_rate"] > 1].shape[0]
```

```
[14]: 110
```



1.2.5 Distribution of Death Rate

As we considered which factor to predict, we created some histograms to observe the distribution of death_rate.

```
[15]: hist_deaths = data.copy()
hist_deaths = hist_deaths[["death_rate"]]

#removed 0% death rates so we can look closely at the distribution of actual_
↳ death rates
hist_deaths = hist_deaths.loc[hist_deaths["death_rate"] != 0]

hist_deaths.hist()
plt.title("Distribution of death_rate in all counties")
plt.xlabel("death_rate")
plt.ylabel("number of counties")

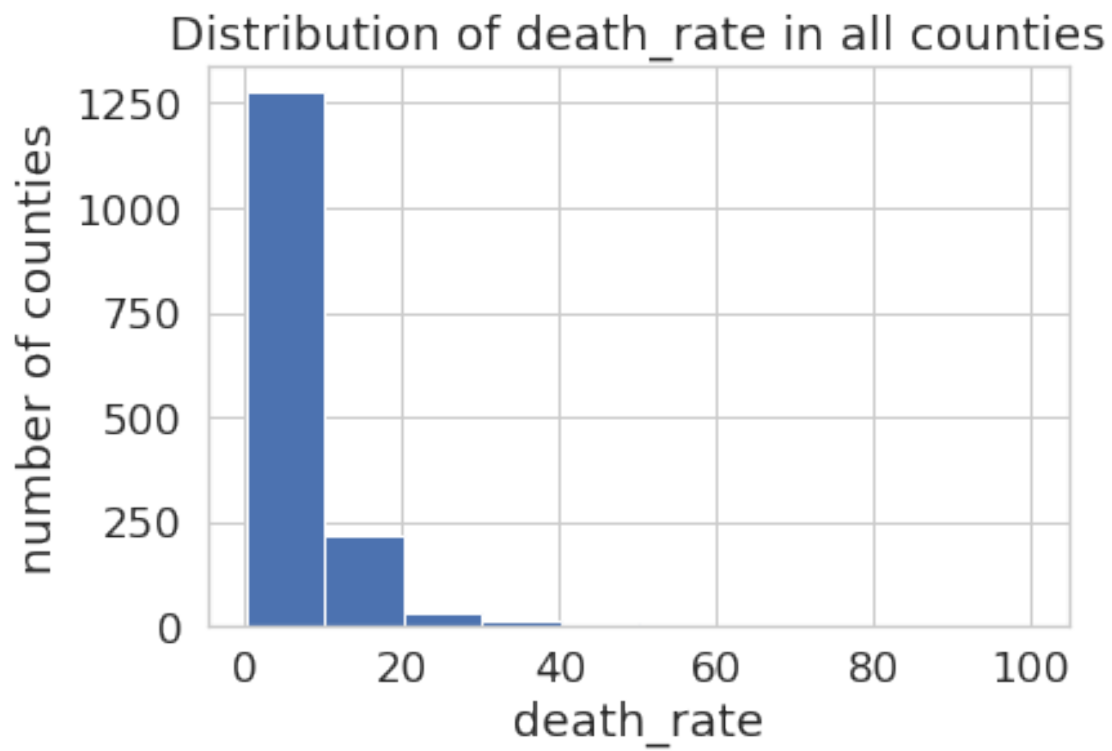
hist_deaths.hist(range = (0,20))
plt.title("Distribution of death_rate within the 0-20% range")
plt.xlabel("death_rate")
plt.ylabel("number of counties")

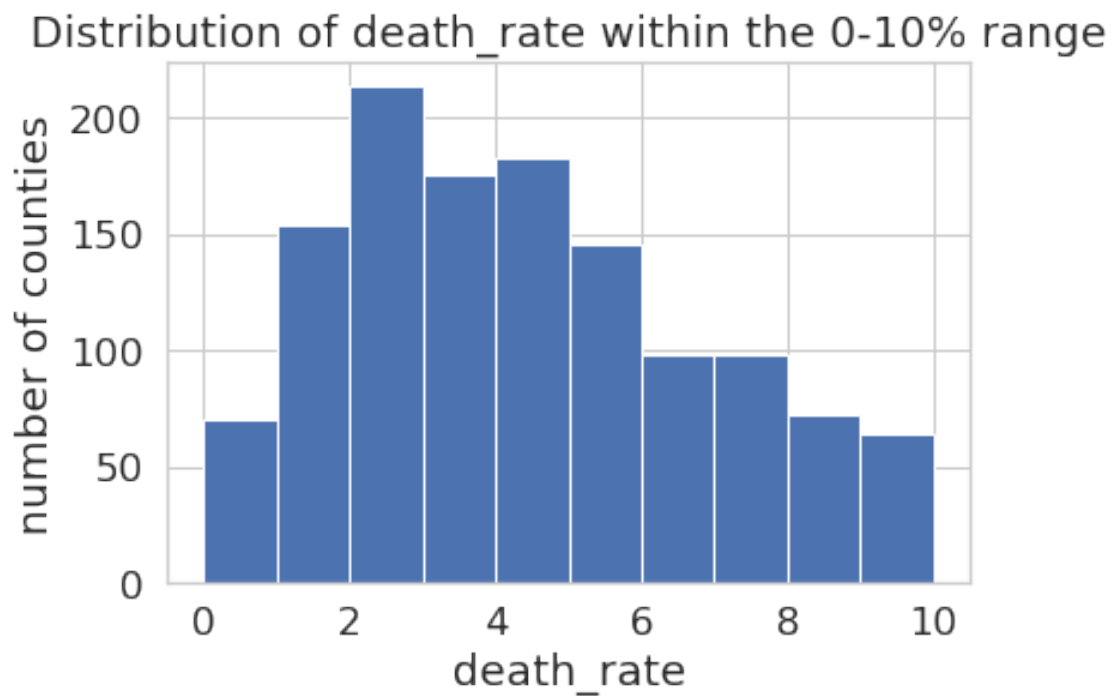
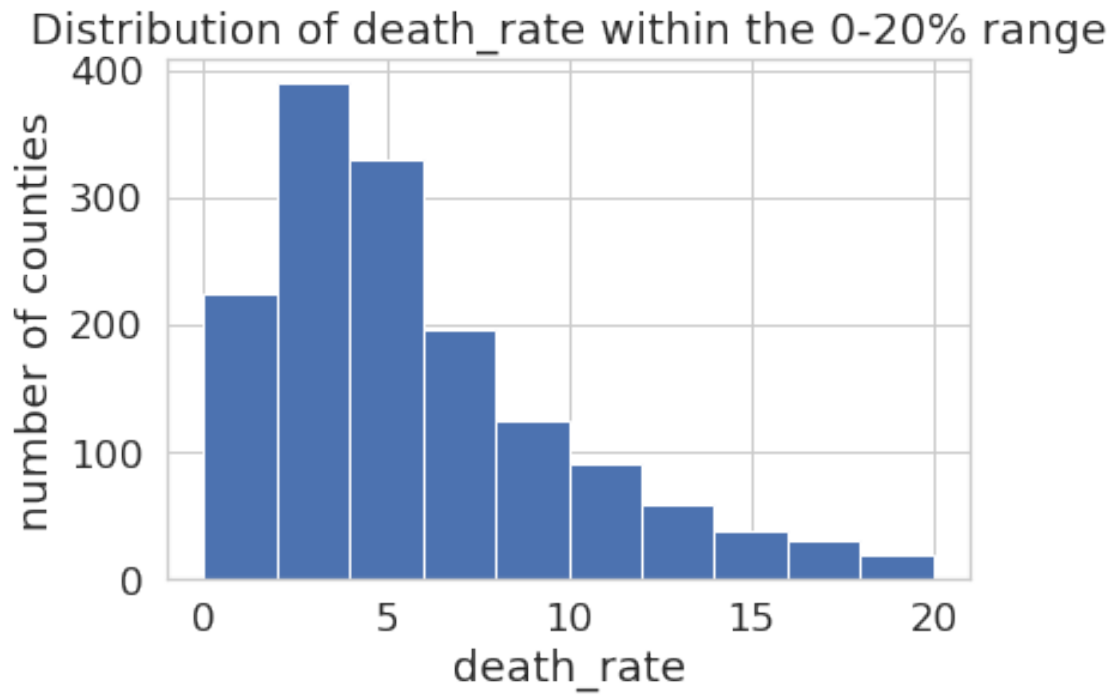
hist_deaths.hist(range = (0, 10))
```

```
plt.title("Distribution of death_rate within the 0-10% range")
plt.xlabel("death_rate")
plt.ylabel("number of counties")

hist_deaths.loc[hist_deaths["death_rate"] > 20].shape[0]
```

[15]: 50





1.2.6 Geographically mapping counties that have any cases and any deaths

We noticed that there were many counties with no cases or no deaths, and we wanted to visually see where these counties were to see if there was a particular state or region that especially lacked cases or deaths.

We created a scatterplot based on latitude and longitude and plotted all counties with cases first and then all counties with deaths.

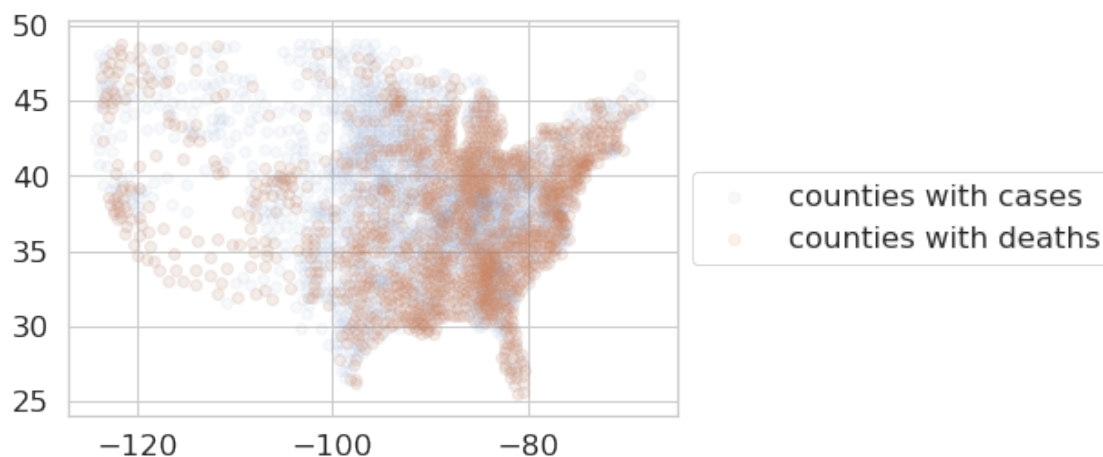
Points where you can just see the blue dot shows there are no deaths in the area. We also played with the alpha value so that the opacity was lower and you could see where cases or deaths were especially concentrated.

```
[16]: maps = data.copy()
      #plot all counties with cases
      plt.scatter(maps["lon"], maps["lat"], alpha = 0.04)

      #remove counties with no deaths and plot the ones with any deaths
      maps = maps.loc[maps["deaths"] != 0]
      plt.scatter(maps["lon"], maps["lat"], alpha = 0.1)

      #add and re-position the legend
      plt.legend(["counties with cases", "counties with deaths"], loc="center left",
        ↪bbox_to_anchor=(1, 0.5))
```

```
[16]: <matplotlib.legend.Legend at 0x7f380c6f1450>
```



1.3 Exploratory Data Analysis on Potential Features

These are features we ended up incorporating into our models.

1.3.1 Median Income and Death Rate

We first made a simple line plot between the median income and the death rate, and saw a general negative correlation between the two.

```
[17]: income = data.copy()

#line plot (hover to show information)
fig = px.line(income, x = "med_income", y = "death_rate",
             ↪hover_name="CountyName")
fig.show()
```

1.3.2 Death Rate and Case Rate for Grouped Median Income

We wanted to get a closer look at that trend, so we did some more thorough investigating. We then grouped the income by the thousands so we would have fewer data points and created a scatterplot with a best fit line to show the correlation even more.

Coupled with the case rate graph, the death rate vs median income graph shows that when income is higher, the death rate of the counties is lower because the case rate stays generally constant throughout all income levels.

```
[18]: income = income.sort_values(["med_income"])

#add column to group income by the thousands
income["income (in k)"] = income["med_income"] // 1000
income = income[["med_income", "income (in k)", "death_rate", "case_rate"]]

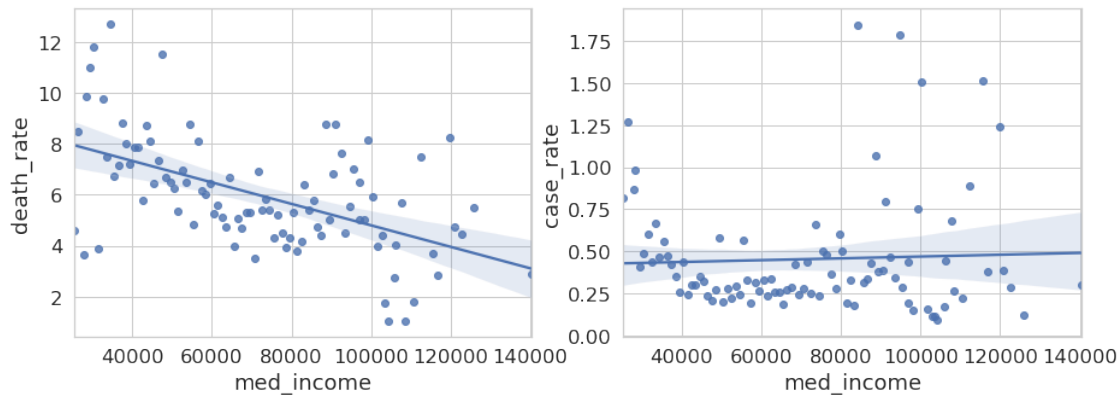
#remove death_rate = 0 since that isn't representative
income = income.loc[income["death_rate"] != 0]

#group by income groups (of thousands)
incomegroup = income.groupby(["income (in k)"]).mean().
    ↪sort_values(["med_income"])

fig, ax = plt.subplots(1,2, figsize=(15,5))

#plot case_rate and death_rate
sns.regplot("med_income", "death_rate", data=incomegroup, ax=ax[0])
sns.regplot("med_income", "case_rate", data=incomegroup, ax=ax[1])
```

```
[18]: <matplotlib.axes._subplots.AxesSubplot at 0x7f3808f28a90>
```



1.3.3 ICU_beds/1000ppl and Death_rate

We then wanted to see if there was a correlation between the number of ICU beds for every 1000 people and death rate, so we created a scatter plot.

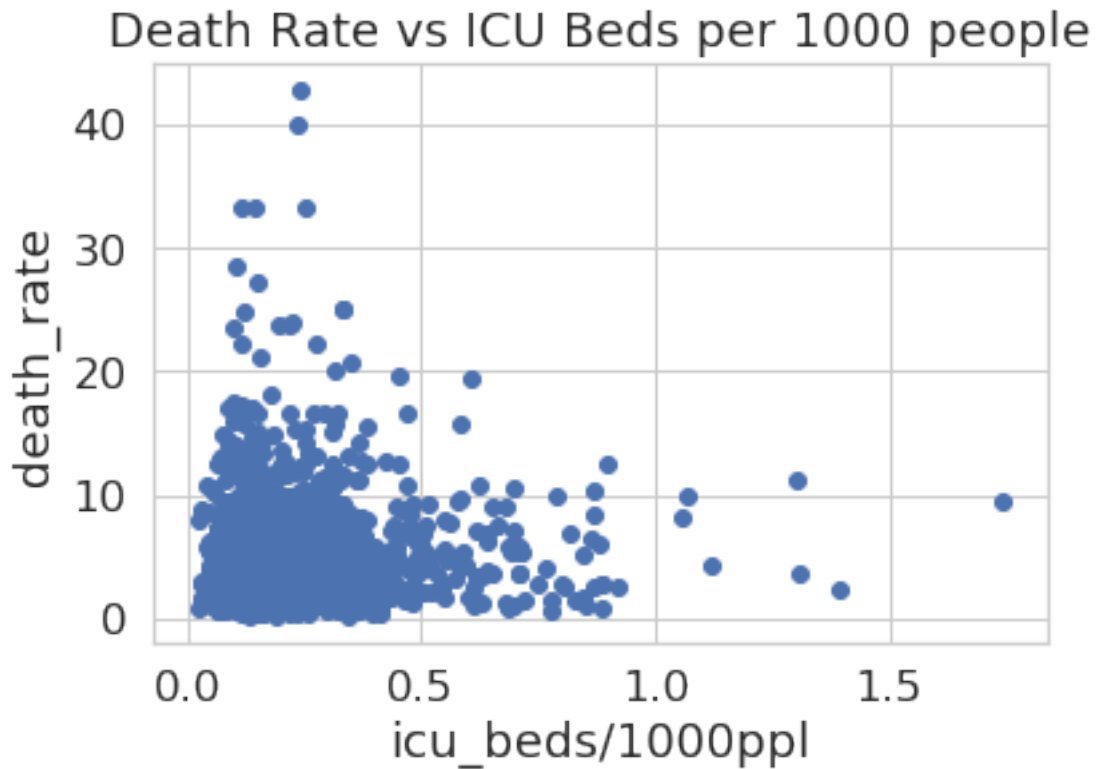
The correlation wasn't that strong, so to use this, we'll have to add on more features.

```
[19]: icu = data.copy()

icu = icu.loc[icu["death_rate"] != 100]
icu = icu.loc[icu["death_rate"] != 0]
icu = icu.loc[icu["icu_beds/1000ppl"] != 0]

plt.scatter("icu_beds/1000ppl", "death_rate", data=icu)
plt.xlabel("icu_beds/1000ppl")
plt.ylabel("death_rate")
plt.title("Death Rate vs ICU Beds per 1000 people")
```

```
[19]: Text(0.5, 1.0, 'Death Rate vs ICU Beds per 1000 people')
```



1.3.4 Age and Death_Rate

We finally wanted to see if there was a correlation between the age factors and death rate, so we created another scatter plot.

This one shows a positive correlation, where counties with higher fractions of 65+ people have a higher death rate. The following shows a positive correlation (but less strong), where counties with higher median ages have a higher death rate.

```
[20]: age = data.copy()

age = age.loc[age["death_rate"] != 100]
age = age.loc[age["death_rate"] != 0]

fig = px.scatter(age, x = "old", y = "death_rate", trendline = "ols")
fig.update_layout(
    title="Age vs Death Rate",
    xaxis_title="% of population with age 65+",
    yaxis_title="death rate",
)
```



```
[21]: fig2 = px.scatter(age, x = "MedianAge2010", y = "death_rate", trendline = "ols")
fig2.update_layout(
    title="Age vs Death Rate",
    xaxis_title="Median Age",
    yaxis_title="death rate",
)
```

1.4 Modeling the Data

1.4.1 Data Cleaning

Before we build our model, we decided we needed to further clean our data so the model could be as accurate as possible. We first removed outliers: * We removed counties where the death rate was 100% (there were 2 counties where this was true, both had 1 case and 1 death). These were large outliers and not representative of the remaining data. * We then decided to remove counties where there were more than 10,000 cases since those were also not representative of the data. There were 20 counties where this was true (out of 2810 total now), so we decided the best course of action was to remove these. * The next decision removed a fairly large number of counties. We decided to remove all counties where the death rate was 0 (of which there were 1265). This was because the goal of the model was to predict death_rate based on various factors, and it didn't make sense for the model to predict a value of 0 since the factors were variable and there can't exist any "perfect" values for the factors that would make them have a 0% death rate.

Then, we cut the dataset down to any factors that we may choose to use later in our models so it's easier to read and process. We counted how many NaN values were in each of the chosen columns, and then seeing that they were a negligible amount (1 or 2 counties), we dropped all counties that had any NaN values in our chosen columns.

```
[22]: print(data.shape[0])
data.columns
```

2812

```
[22]: Index(['countyFIPS', 'STATEFP', 'COUNTYFP', 'CountyName', 'StateName', 'State',
        'lat', 'lon', 'POP_LATITUDE', 'POP_LONGITUDE',
        ...,
        'FIPStxt_y', 'unemploy_rate', 'med_income', 'case_rate', 'death_rate',
        'old', 'inMedicare', 'medicare_rate', 'hospitals/1000ppl',
        'icu_beds/1000ppl'],
        dtype='object', length=101)
```

```
[23]: traintest = data.copy()

#remove outliers
traintest.loc[traintest["death_rate"] == 100].shape[0] #there are 2 counties_
↳where this is true
traintest = traintest.loc[traintest["death_rate"] != 100]
```

```

traintest.loc[traintest["confirmed"] > 10000].shape[0] #there are 20 of these
↳ outliers with significantly more cases
traintest = traintest.loc[traintest["confirmed"] < 10000]

traintest = traintest.loc[traintest["deaths"] > 0.00]

traintest = traintest.rename(columns={"PopulationDensityperSqMile2010":
↳ "density"})

traintest = traintest[["confirmed",
                        "deaths",
                        "death_rate",
                        "case_rate",
                        "old",
                        "MedianAge2010",
                        "density",
                        "inMedicare",
                        "medicare_rate",
                        "pov_pct",
                        "unemploy_rate",
                        "med_income",
                        "lat",
                        "lon",
                        "FracMale2017",
                        "DiabetesPercentage",
                        "HeartDiseaseMortality",
                        "StrokeMortality",
                        "Smokers_Percentage",
                        "RespMortalityRate2014",
                        "hospitals/1000ppl",
                        "icu_beds/1000ppl",
                        "SVIPercentile",
                        "TotalM.D.'s,TotNon-FedandFed2017",
                        "Rural-UrbanContinuumCode2013",
                        "#FTEHospitalTotal2017"]]
#print(traintest.isna().sum())
traintest = traintest.dropna()

traintest = (traintest.loc[traintest["confirmed"] > 5])
traintest.loc[traintest["confirmed"] > 5].shape[0]

```

[23]: 1489

1.4.2 Train-Test Split

Next we split our cleaned dataset into a training data set and a testing data set using sklearn. We put 85% of our data into the training set and 15% into the testing set. We picked a `random_state` so that the split was pseudorandom.

```
[24]: train, test = train_test_split(traintest, test_size=0.15, random_state=42)
      train_c = train
      test_c = test
```

```
[25]: # from sklearn import linear_model as lm
      train_c
```

```
[25]: confirmed deaths death_rate case_rate old MedianAge2010 \
1765      1558      79    5.070603    0.586974  17.595289      39.9
2116       213     13    6.103286    0.091742  13.942620      37.8
2426        22      1    4.545455    0.078515  20.089222      42.0
2072        39      1    2.564103    0.066662  16.747573      38.4
813         6      1   16.666667    0.038069  21.686441      43.3
...      ...      ...      ...      ...      ...
2240     3453     172    4.981176    0.821845  16.849140      39.1
2617      203      2    0.985222    0.091182  10.239365      30.4
1757       44      2    4.545455    0.057522  19.179533      40.7
3047       34      1    2.941176    0.059277  17.280937      41.0
2236     1526     127    8.322412    0.125241  18.515707      41.3

      density inMedicare medicare_rate pov_pct ... HeartDiseaseMortality \
1765    494.1   77.058475     22.260190    12.8 ...                203.7
2116    530.0   80.931568     16.949861     5.2 ...                155.9
2426    47.6   73.813112     28.415418    18.6 ...                315.7
2072   121.3   77.136540     22.420689    11.1 ...                174.5
813     32.6   81.084427     25.626547    11.4 ...                162.4
...      ...      ...      ...      ...      ...
2240    480.4   79.603396     20.716074    11.4 ...                171.2
2617    231.7   75.262671     13.594693    13.2 ...                131.2
1757    109.1   75.991409     24.955225    10.2 ...                148.0
3047     74.2   78.829787     22.943617     7.6 ...                163.2
2236   1675.7   82.062767     22.006201    11.7 ...                185.8

      StrokeMortality Smokers_Percentage RespMortalityRate2014 \
1765                31.8         15.644195         50.72
2116                36.8         15.228085         57.85
2426                57.6         22.895309         61.91
2072                35.1         19.307126         73.64
813                 30.4         15.947325         51.87
...      ...      ...      ...
2240                50.8         16.583595         47.62
```

| | | | |
|------|------|-----------|-------|
| 2617 | 36.2 | 14.147185 | 43.77 |
| 1757 | 30.9 | 16.624635 | 64.46 |
| 3047 | 32.1 | 14.981969 | 58.82 |
| 2236 | 35.7 | 16.991039 | 48.80 |

| | hospitals/1000ppl | icu_beds/1000ppl | SVIPercentile \ |
|------|-------------------|------------------|-----------------|
| 1765 | 0.007535 | 0.241119 | 0.7911 |
| 2116 | 0.004307 | 0.245507 | 0.0478 |
| 2426 | 0.035689 | 0.214133 | 0.7443 |
| 2072 | 0.034186 | 0.205114 | 0.4325 |
| 813 | 0.063448 | 0.000000 | 0.3596 |
| ... | ... | ... | ... |
| 2240 | 0.007140 | 0.168986 | 0.6162 |
| 2617 | 0.013475 | 0.188653 | 0.4924 |
| 1757 | 0.013073 | 0.352973 | 0.1417 |
| 3047 | 0.034869 | 0.174344 | 0.1204 |
| 2236 | 0.012311 | 0.476014 | 0.2506 |

| | TotalM.D.'s,TotNon-FedandFed2017 | Rural-UrbanContinuumCode2013 \ |
|------|----------------------------------|--------------------------------|
| 1765 | 779.0 | 2.0 |
| 2116 | 651.0 | 1.0 |
| 2426 | 31.0 | 6.0 |
| 2072 | 60.0 | 4.0 |
| 813 | 9.0 | 7.0 |
| ... | ... | ... |
| 2240 | 890.0 | 2.0 |
| 2617 | 296.0 | 1.0 |
| 1757 | 215.0 | 4.0 |
| 3047 | 59.0 | 2.0 |
| 2236 | 8995.0 | 1.0 |

| | #FTEHospitalTotal2017 |
|------|-----------------------|
| 1765 | 6125.0 |
| 2116 | 1455.0 |
| 2426 | 283.0 |
| 2072 | 937.0 |
| 813 | 173.0 |
| ... | ... |
| 2240 | 8214.0 |
| 2617 | 1399.0 |
| 1757 | 1260.0 |
| 3047 | 806.0 |
| 2236 | 37225.0 |

[1265 rows x 26 columns]

[26]: test_c

[26]:

| | confirmed | deaths | death_rate | case_rate | old | MedianAge2010 | \ |
|------|-----------|--------|------------|-----------|-----------|---------------|------|
| 1896 | 35 | 3 | 8.571429 | 0.050342 | 24.017605 | | 45.8 |
| 483 | 50 | 1 | 2.000000 | 0.125247 | 14.155457 | | 36.2 |
| 447 | 2060 | 35 | 1.699029 | 1.019055 | 14.447336 | | 34.5 |
| 1906 | 399 | 11 | 2.756892 | 0.120061 | 11.769326 | | 31.0 |
| 2111 | 413 | 36 | 8.716707 | 0.207927 | 21.151203 | | 42.8 |
| ... | ... | ... | ... | ... | ... | | ... |
| 380 | 20 | 1 | 5.000000 | 0.241051 | 14.318428 | | 33.4 |
| 1846 | 1890 | 156 | 8.253968 | 0.254554 | 16.835068 | | 38.5 |
| 38 | 99 | 2 | 2.020202 | 0.107158 | 19.603407 | | 40.4 |
| 96 | 22 | 1 | 4.545455 | 0.104275 | 38.406484 | | 53.9 |
| 2341 | 59 | 1 | 1.694915 | 0.223502 | 17.838473 | | 38.8 |

| | density | inMedicare | medicare_rate | pov_pct | ... | HeartDiseaseMortality | \ |
|------|---------|------------|---------------|---------|-----|-----------------------|-------|
| 1896 | 131.3 | 79.446079 | 28.355676 | 9.8 | ... | | 164.1 |
| 483 | 115.0 | 68.404160 | 19.511034 | 15.2 | ... | | 210.5 |
| 447 | 457.5 | 78.287274 | 18.490908 | 13.2 | ... | | 145.8 |
| 1906 | 489.7 | 69.597926 | 16.479403 | 17.0 | ... | | 188.2 |
| 2111 | 340.2 | 79.151284 | 26.456625 | 17.6 | ... | | 213.2 |
| ... | ... | ... | ... | ... | ... | | ... |
| 380 | 24.7 | 71.045392 | 17.524406 | 26.1 | ... | | 219.3 |
| 1846 | 1132.6 | 78.660949 | 21.450179 | 14.4 | ... | | 149.3 |
| 38 | 138.9 | 75.752346 | 24.457987 | 14.0 | ... | | 219.7 |
| 96 | 4.6 | 77.438727 | 29.201820 | 23.7 | ... | | 182.7 |
| 2341 | 60.3 | 70.396666 | 24.543526 | 30.0 | ... | | 368.7 |

| | StrokeMortality | Smokers_Percentage | RespMortalityRate2014 | \ |
|------|-----------------|--------------------|-----------------------|---|
| 1896 | 40.5 | 15.646474 | 52.63 | |
| 483 | 48.5 | 17.836136 | 98.05 | |
| 447 | 39.9 | 15.636843 | 61.37 | |
| 1906 | 42.5 | 18.090857 | 68.59 | |
| 2111 | 43.0 | 19.676253 | 63.93 | |
| ... | ... | ... | ... | |
| 380 | 49.7 | 21.822201 | 116.44 | |
| 1846 | 32.3 | 15.151442 | 37.67 | |
| 38 | 49.9 | 16.656004 | 75.94 | |
| 96 | 31.4 | 15.521010 | 45.08 | |
| 2341 | 63.8 | 21.379465 | 92.22 | |

| | hospitals/1000ppl | icu_beds/1000ppl | SVIPercentile | \ |
|------|-------------------|------------------|---------------|---|
| 1896 | 0.014384 | 0.115068 | 0.4255 | |
| 483 | 0.025049 | 0.000000 | 0.7134 | |
| 447 | 0.004947 | 0.420484 | 0.6618 | |
| 1906 | 0.003009 | 0.144435 | 0.8640 | |
| 2111 | 0.010069 | 0.211452 | 0.4640 | |
| ... | ... | ... | ... | |
| 380 | 0.000000 | 0.000000 | 0.9615 | |

| | | | |
|------|----------|----------|--------|
| 1846 | 0.006734 | 0.257248 | 0.5204 |
| 38 | 0.010824 | 0.519554 | 0.4468 |
| 96 | 0.094796 | 0.142194 | 0.9236 |
| 2341 | 0.000000 | 0.000000 | 0.9869 |

| | TotalM.D.'s,TotNon-FedandFed2017 | Rural-UrbanContinuumCode2013 | \ |
|------|----------------------------------|------------------------------|-----|
| 1896 | 212.0 | | 4.0 |
| 483 | 13.0 | | 3.0 |
| 447 | 475.0 | | 3.0 |
| 1906 | 851.0 | | 2.0 |
| 2111 | 290.0 | | 2.0 |
| ... | ... | ... | |
| 380 | 1.0 | | 9.0 |
| 1846 | 4297.0 | | 1.0 |
| 38 | 195.0 | | 3.0 |
| 96 | 16.0 | | 6.0 |
| 2341 | 15.0 | | 6.0 |

| | #FTEHospitalTotal2017 |
|------|-----------------------|
| 1896 | 903.0 |
| 483 | 214.0 |
| 447 | 4546.0 |
| 1906 | 8580.0 |
| 2111 | 2707.0 |
| ... | ... |
| 380 | 0.0 |
| 1846 | 23546.0 |
| 38 | 1250.0 |
| 96 | 947.0 |
| 2341 | 0.0 |

[224 rows x 26 columns]

1.4.3 Defining Functions and Setting Up for the Models

Here, we defined a function to calculate the root mean squared error (RMSE) given actual and predicted values. This was created while referencing the function defined in lecture.

We also initialized a dictionary to hold the models so we can compare them later, and arrays to hold the RMSEs for training data, cross validation, and testing data so we could compare them at the end. This was also created in reference to the Cross Validation lecture.

Since we are predicting death_rate that variable is set (and this structure makes it easy to switch to case_rate or any other factor we want to predict!).

We also find the range of the death_rates in our data so we can contextualize our RMSE later.

```
[27]: def rmse(actual, predicted):
        return np.sqrt(np.mean((actual - predicted)**2))
```

```
[28]: def standardize(data):
        return (data - np.mean(data)) / np.std(data)
```

```
[114]: models = {}
        training_rmse = []
        validate_rmse = []
        test_rmse = []
```

```
[30]: predicting = "death_rate"

        print(min(train_c[predicting]), " - ", max(train_c[predicting]))
        print(np.mean(train_c[predicting]))

        standard = standardize(train_c[predicting])
        print(min(standard), " - ", max(standard))
```

```
0.15384615384615385 - 42.857142857142854
6.01213217346454
-1.2663776434065634 - 7.964735358538478
```

1.5 Feature Engineering

We started the process of evaluating what features are best by creating different pairplots to compare death_rate and case_rate to certain factors. Below is the pairplot we created for some economic factors.

(Note: we created several different pairplots but including all of them caused my kernel to crash, so we're only showing one below)

```
[31]: #sns.pairplot(train_c[["death_rate", "case_rate", "pov_pct", "unemploy_rate",
        ↪ "med_income"]])
```

1.5.1 General Health Features Model

In terms of general health, the factors we took into account were: * the percentage of “old” people in that county (over 65 years old) * the median age of people in the county * the percentage of the population who has diabetes * the number of heart disease mortalities (per 100,000 people) per year of that county (from 2014-2016) * the number of stroke mortalities (per 100,000 people) per year of that county (from 2014-2016) * the percentage of the population who are smokers * the respiratory mortality rate (per 100,000 people) from the year 2014.

```
[115]: def health_cols(data):
        return data[[predicting,
```

```

        "old",
        "MedianAge2010",
        "DiabetesPercentage",
        "HeartDiseaseMortality",
        "StrokeMortality",
        "Smokers_Percentage",
        "RespMortalityRate2014"
    ]]

```

We create the `X_train` and `y_train` variables for the model, and then create the Linear Regression model that we're using.

```

[116]: X_train_h = health_cols(train_c).drop([predicting], axis = 1)
y_train = train_c[predicting]

regr = lm.LinearRegression()

X_train_h

```

```

[116]:
      old  MedianAge2010  DiabetesPercentage  HeartDiseaseMortality \
1765  17.595289         39.9                10.4                 203.7
2116  13.942620         37.8                 9.5                 155.9
2426  20.089222         42.0                 9.9                 315.7
2072  16.747573         38.4                11.1                 174.5
813   21.686441         43.3                 6.3                 162.4
...
2240  16.849140         39.1                11.2                 171.2
2617  10.239365         30.4                 8.5                 131.2
1757  19.179533         40.7                 7.1                 148.0
3047  17.280937         41.0                 9.4                 163.2
2236  18.515707         41.3                 8.1                 185.8

      StrokeMortality  Smokers_Percentage  RespMortalityRate2014
1765              31.8        15.644195              50.72
2116              36.8        15.228085              57.85
2426              57.6        22.895309              61.91
2072              35.1        19.307126              73.64
813               30.4        15.947325              51.87
...
2240              50.8        16.583595              47.62
2617              36.2        14.147185              43.77
1757              30.9        16.624635              64.46
3047              32.1        14.981969              58.82
2236              35.7        16.991039              48.80

[1265 rows x 7 columns]

```

In the following cells, we fit the model and use it to predict values based on `X_train`. We also

use the previously defined `rmse` function to calculate the training error. We also use the score function built into models, although we mostly used the RMSE to determine and score the model's performance.

```
[117]: regr.fit(X_train_h, y_train)
      y_fitted = regr.predict(X_train_h)
```

```
[118]: models["health"] = regr
```

```
[119]: training_error = rmse(y_fitted, y_train)
      training_rmse.append(training_error)
      training_error
```

```
[119]: 4.429037715031477
```

```
[120]: print(regr.score(X_train_h, y_train))
```

```
0.08334885794761948
```

Next, we perform cross validation on this model using the `cross_val_score` method from sklearn. Instead of the default scoring, we used one of the built-in scoring methods which was “`neg_root_mean_squared_error`”. Since we didn’t actually want the value to be negative we then make them positive again. We then take the mean of these 5 cross validated RMSEs and will use that later to compare.

```
[121]: c_scores = cross_val_score(regr, X_train_h, y_train, cv=5,
      ↪scoring="neg_root_mean_squared_error")
      c_scores = c_scores * -1
      c_scores.mean()
      validate_rmse.append(c_scores.mean())
      c_scores
```

```
[121]: array([4.16004279, 5.03249353, 4.18326695, 4.8132614 , 4.14195952])
```

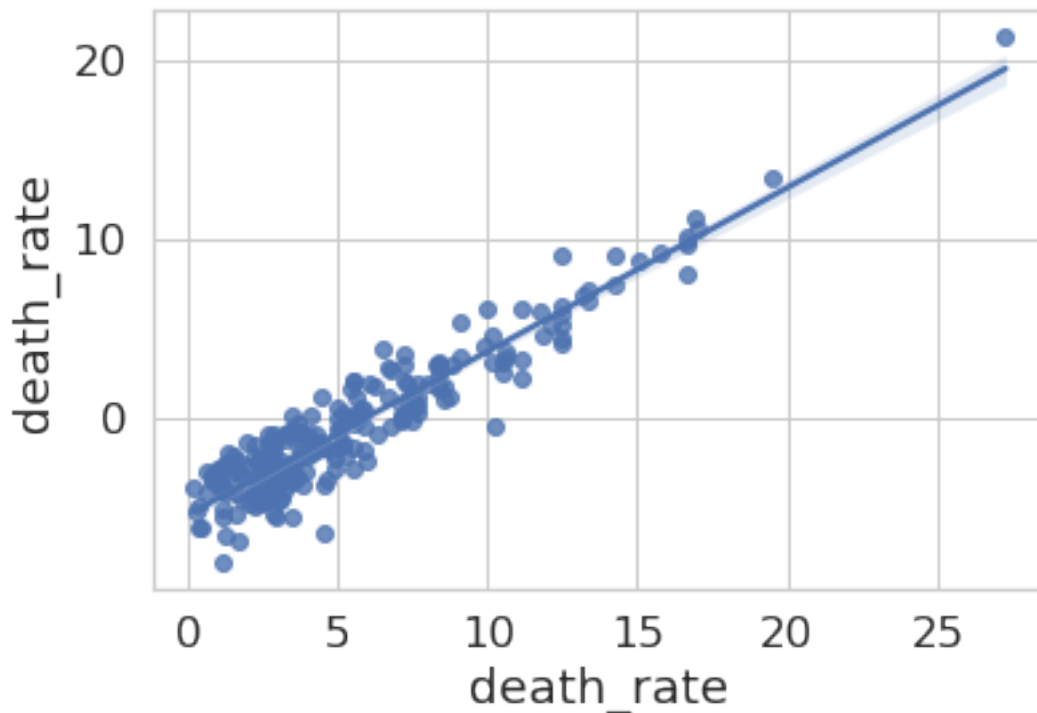
Finally, we predict our values for our testing data set, and calculate the testing RMSE. We also find the residuals and plot them against the actual values (this was referenced from part homeworks and labs).

```
[122]: X_test_h = health_cols(test_c).drop([predicting], axis = 1)
      y_test = test_c[predicting]

      y_predicted = regr.predict(X_test_h)
      residuals = y_test - y_predicted
      ax = sns.regplot(y_test, residuals)

      test_error = rmse(y_predicted, y_test)
      test_rmse.append(test_error)
      test_error
```

[122]: 4.165044006828287



1.5.2 Healthcare System Features Model

In terms of the healthcare system, the factors we took into account were: * the number of hospitals per 1,000 people in the county * the number of ICU beds per 1,000 people in the county * the percent of people enrolled in Medicare who are eligible for Medicare * the number of MDs in each county in 2017 * the number of full-time employees at hospitals in 2017 in the county * the rural-urban continuum code that determines how rural or urban a county is on a scale from 1 to 9

```
[123]: def healthcare_cols(data):  
        return data[[predicting,  
                      "hospitals/1000ppl",  
                      "icu_beds/1000ppl",  
                      "inMedicare",  
                      "TotalM.D.'s,TotNon-FedandFed2017",  
                      "#FTEHospitalTotal2017",  
                      "Rural-UrbanContinuumCode2013"  
                      ]]
```

The model training, error finding, cross validation, and testing process mirrors the one above for the general health features.

```
[124]: X_train_hc = healthcare_cols(train_c).drop([predicting], axis = 1)
y_train = train_c[predicting]

regr2 = lm.LinearRegression()
```

```
[125]: regr2.fit(X_train_hc, y_train)
y_fitted = regr2.predict(X_train_hc)
```

```
[126]: models["healthcare"] = regr2
```

```
[127]: training_error = rmse(y_fitted, y_train)
training_rmse.append(training_error)
training_error
```

```
[127]: 4.494261786979374
```

```
[128]: print(regr2.score(X_train_hc, y_train))
```

```
0.056152001736684325
```

```
[129]: c_scores = cross_val_score(regr2, X_train_hc, y_train, cv=5,
    ↳scoring="neg_root_mean_squared_error")
c_scores = c_scores * -1
c_scores.mean()
validate_rmse.append(c_scores.mean())
c_scores
```

```
[129]: array([4.27377157, 5.04418444, 4.28248583, 4.84825898, 4.12799525])
```

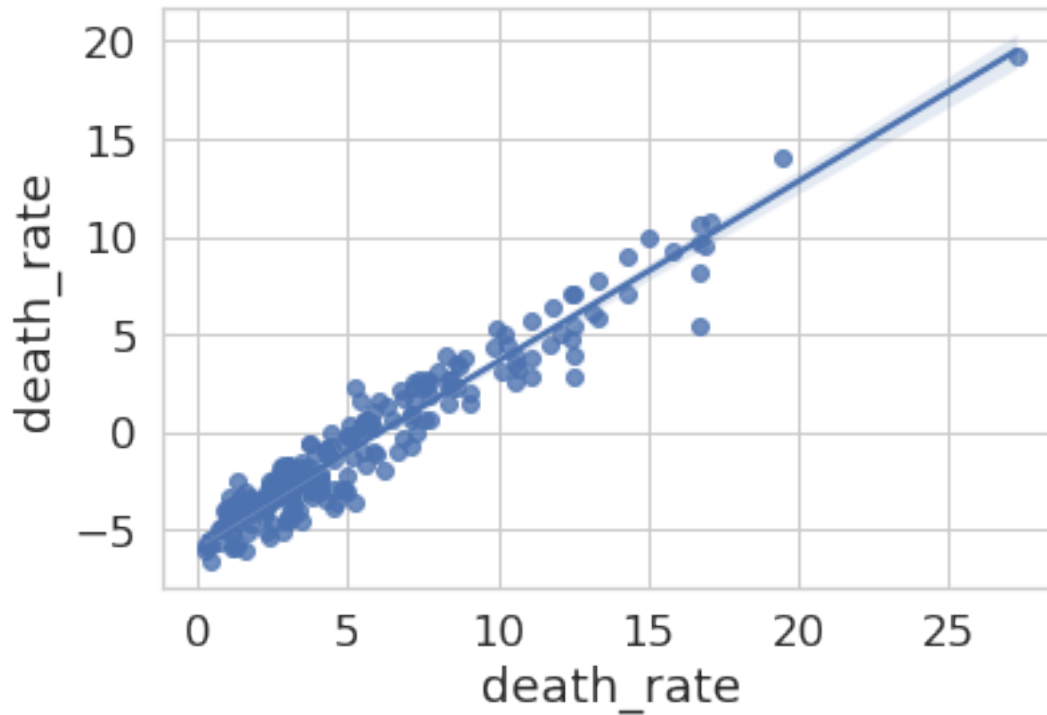
```
[130]: X_test_hc = healthcare_cols(test_c).drop([predicting], axis = 1)
y_test = test_c[predicting]

y_predicted = regr2.predict(X_test_hc)
residuals = y_test - y_predicted
ax = sns.regplot(y_test, residuals)

test_error = rmse(y_predicted, y_test)
test_rmse.append(test_error)

test_error
```

```
[130]: 4.101943608331466
```



1.5.3 Economic State Features Model

In terms of the economic state, the factors we took into account were: * median income of the county * the unemployment rate of the county * the percentage of poverty in the county * the percentage of people eligible for Medicare in the county * the county's overall percentile ranking indicating the CDC's Social Vulnerability Index (SVI)

```
[131]: def econ_cols(data):
        return data[[predicting,
                      "SVIPercentile",
                      "medicare_rate",
                      "pov_pct",
                      "unemploy_rate",
                      "med_income"
                      ]]
```

The model training, error finding, cross validation, and testing process mirrors the one above for the general health features.

```
[132]: X_train_e = econ_cols(train_c).drop([predicting], axis = 1)
        y_train = train_c[predicting]

        regr3 = lm.LinearRegression()
```

```
[133]: regr3.fit(X_train_e, y_train)
y_fitted = regr3.predict(X_train_e)
```

```
[134]: models["economic"] = regr3
```

```
[135]: training_error = rmse(y_fitted, y_train)
training_rmse.append(training_error)
training_error
```

```
[135]: 4.4769016717887125
```

```
[136]: print(regr3.score(X_train_e, y_train))
```

```
0.06342957689439865
```

```
[137]: c_scores = cross_val_score(regr3, X_train_e, y_train, cv=5,
    ↪scoring="neg_root_mean_squared_error")
c_scores = c_scores * -1
c_scores.mean()
validate_rmse.append(c_scores.mean())
c_scores
```

```
[137]: array([4.15907729, 5.1289684 , 4.23665803, 4.81163882, 4.10209386])
```

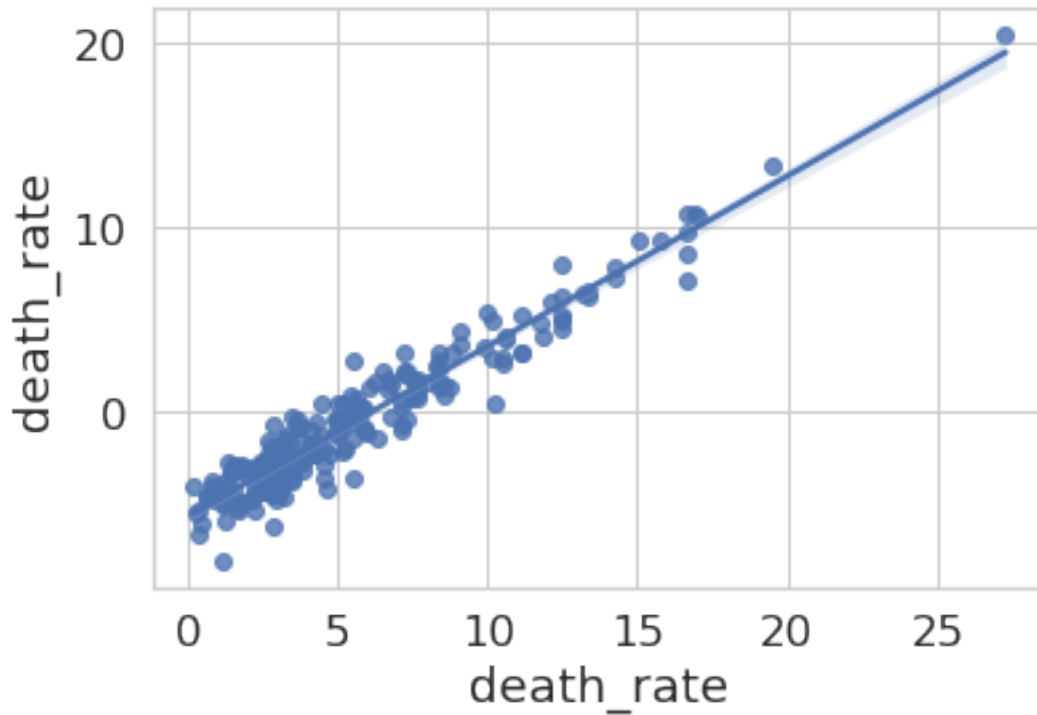
```
[138]: X_test_e = econ_cols(test_c).drop([predicting], axis = 1)
y_test = test_c[predicting]

y_predicted = regr3.predict(X_test_e)
residuals = y_test - y_predicted
ax = sns.regplot(y_test, residuals)

test_error = rmse(y_predicted, y_test)
test_rmse.append(test_error)

test_error
```

```
[138]: 4.114396933982724
```



1.5.4 All Features Model

Here we combined all the features from the previous 3 models.

```
[139]: def all_cols(data):
        return data[[predicting,
                       "old",
                       "DiabetesPercentage",
                       "HeartDiseaseMortality",
                       "StrokeMortality",
                       "Smokers_Percentage",
                       "RespMortalityRate2014",
                       "hospitals/1000ppl",
                       "icu_beds/1000ppl",
                       "inMedicare",
                       "TotalM.D.'s,TotNon-FedandFed2017",
                       "#FTEHospitalTotal2017",
                       "Rural-UrbanContinuumCode2013",
                       "SVIPercentile",
                       "medicare_rate",
                       "pov_pct",
                       "unemploy_rate",
                       "med_income"]]
```

```
]]
```

The model training, error finding, cross validation, and testing process mirrors the one above for the general health features.

```
[140]: X_train_a = all_cols(train_c).drop([predicting], axis = 1)
       y_train = train_c[predicting]

       regr4 = lm.LinearRegression()
```

```
[141]: regr4.fit(X_train_a, y_train)
       y_fitted = regr4.predict(X_train_a)
```

```
[142]: models["all"] = regr4
```

```
[143]: training_error = rmse(y_fitted, y_train)
       training_rmse.append(training_error)
       training_error
```

```
[143]: 4.368192572487627
```

```
[144]: print(regr4.score(X_train_a, y_train))
```

```
0.10836136364220872
```

```
[145]: c_scores = cross_val_score(regr4, X_train_a, y_train, cv=5,
    ↪scoring="neg_root_mean_squared_error")
       c_scores = c_scores * -1
       c_scores.mean()
       validate_rmse.append(c_scores.mean())
       c_scores
```

```
[145]: array([4.23722201, 4.94783012, 4.21891976, 4.72464312, 4.10778179])
```

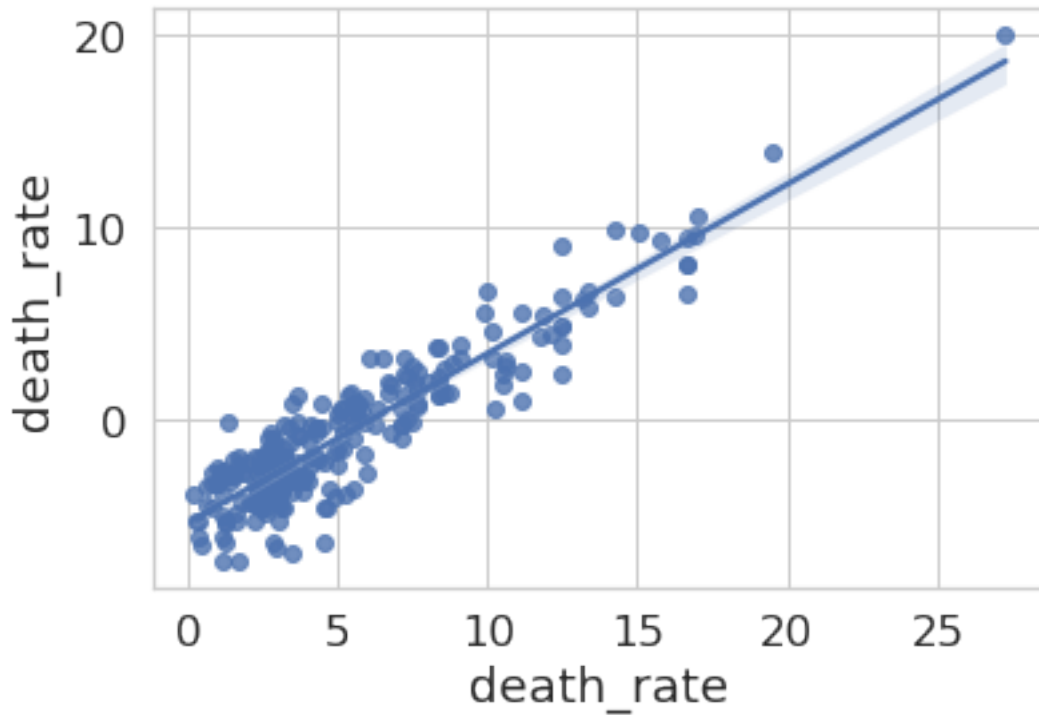
```
[146]: X_test_a = all_cols(test_c).drop([predicting], axis = 1)
       y_test = test_c[predicting]

       y_predicted = regr4.predict(X_test_a)
       residuals = y_test - y_predicted
       ax = sns.regplot(y_test, residuals)

       test_error = rmse(y_predicted, y_test)
       test_rmse.append(test_error)

       test_error
```

```
[146]: 4.096416655290851
```



1.5.5 Comparing Models

Here we compare the previous 4 models based on the training rmse, the cross validation rmse, and the testing rmse. This process was created by referencing lecture code.

[147]: *#adapted from lecture code*

```
def compare_models(models):
    names = list(models.keys())
    fig = go.Figure([
        go.Bar(x = names, y = training_rmse, name="Training RMSE"),
        go.Bar(x = names, y = validate_rmse, name="CV RMSE"),
        go.Bar(x = names, y = test_rmse, name="Test RMSE", opacity=.3)])
    return fig

training_rmse
```

[147]: [4.429037715031477, 4.494261786979374, 4.4769016717887125, 4.368192572487627]

```
[150]: fig = compare_models(models)
fig.update_yaxes(range=[4.05,4.52], title="RMSE")
```

[]:

1.6 Appendix

As a note, we did not think of making an appendix at the end until we saw this on Piazza, so not all our ideas and original failed tests and models and features are in this appendix. We only added old models from the last day of us working on it, so this is not all-encompassing

1.6.1 Baseline Features Model

```
[66]: def no_cols(data):  
      return data[[predicting,  
                  "lat"  
                  ]]
```

```
[67]: X_train_n = no_cols(train_c).drop([predicting], axis = 1)  
      y_train = train_c[predicting]  
  
      regr0 = lm.LinearRegression()
```

```
[68]: regr0.fit(X_train_n, y_train)  
      y_fitted = regr0.predict(X_train_n)
```

```
[69]: models["none"] = regr0
```

```
[70]: training_error = rmse(y_fitted, y_train)  
      training_rmse.append(training_error)  
      training_error
```

```
[70]: 4.6197655238922515
```

```
[71]: print(regr0.score(X_train_n, y_train))
```

```
0.002701435567487143
```

```
[72]: c_scores = cross_val_score(regr0, X_train_n, y_train, cv=5,  
    ↪scoring="neg_root_mean_squared_error")  
      c_scores = c_scores * -1  
      print(c_scores.mean())  
      validate_rmse.append(c_scores.mean())  
      c_scores
```

```
4.632113297626638
```

```
[72]: array([4.38060417, 5.2380538 , 4.29117572, 4.94626447, 4.30446833])
```

```
[test data below]
```

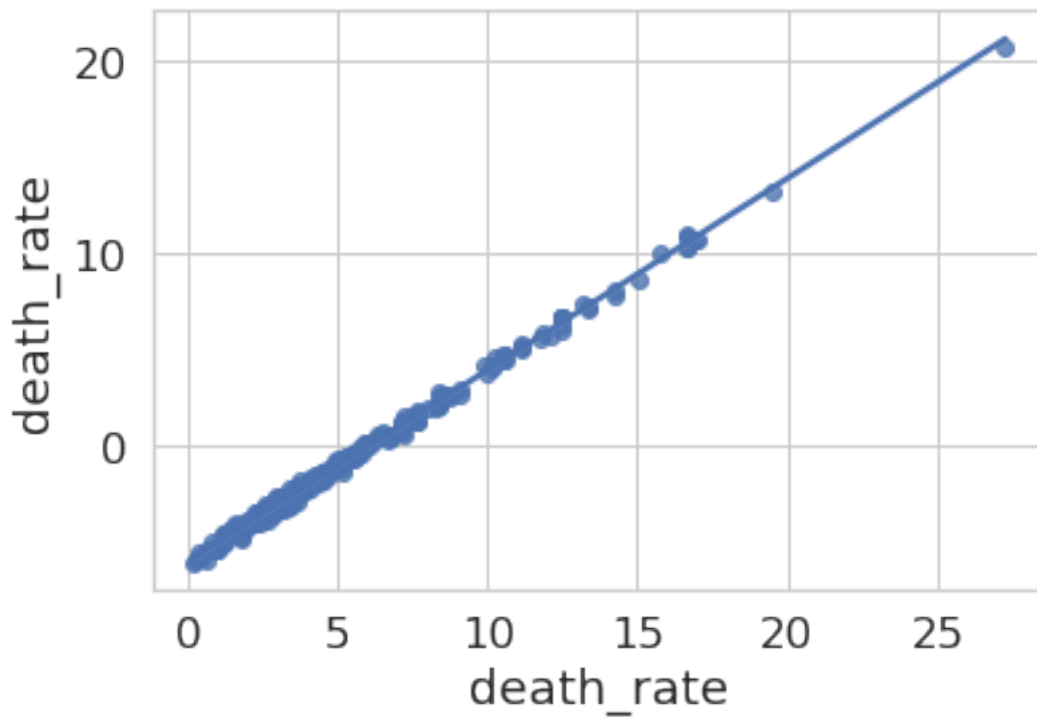
```
[73]: X_test_n = no_cols(test_c).drop([predicting], axis = 1)
y_test = test_c[predicting]

y_predicted = regr0.predict(X_test_n)
residuals = y_test - y_predicted
ax = sns.regplot(y_test, residuals)

test_error = rmse(y_predicted, y_test)
test_rmse.append(test_error)

test_error
```

[73]: 4.274476033190996



1.6.2 Attempt at Standardizing Model

```
[74]: stda = preprocessing.scale(X_train_h)
stda
```

[74]: array([[0.04333002, 0.23610179, -0.10396778, ..., -1.06576233,
 -0.57728857, -0.77990472],
 [-0.84730103, -0.25744905, -0.34697191, ..., -0.48510175,

```

-0.70850521, -0.3204902 ],
[ 0.65142622, 0.72965263, -0.23897007, ..., 1.93044627,
 1.70929051, -0.0588881 ],
...,
[ 0.4296165 , 0.42412116, -0.99498293, ..., -1.17028124,
-0.26811468, 0.10541863],
[-0.03331847, 0.49462842, -0.37397237, ..., -1.0309227 ,
-0.78611586, -0.25798921],
[ 0.26775567, 0.56513569, -0.72497834, ..., -0.61284708,
-0.15257239, -0.90361803]])

```